



Technical University of Denmark



High Performance Computing

FORTRAN, OpenMP and MPI

41391

Content

- Day 3:
 - Array features
 - Zero size array.
 - Assumed size /Assumed shape arrays.
 - Automatic/ALLOCATABLE/POINTER arrays.
 - ALLOCATE/DEALLOCATE/NULLIFY.
 - FORALL, WHERE.
 - PURE/ELEMENTAL procedures.

Zero size arrays

- FORTRAN allows zero size arrays:

Example: solve lower-triangular set of linear equations: $a * x = b$:


DO i=1,n

$x(i) = b(i)/a(i,i)$

$b(i+1:n) = b(i+1:n) - a(i+1:n,i) * x(i)$

ENDDO

- The subsection of b has zero length when $i = n$



Zero size arrays

- Double DO version (preferred solution):

```
DO i=1,n
```

```
    x(i) = b(i)/a(i,i)
```

```
    DO j=i+1,n
```

```
        b(j) = b(j) - a(j,i)*x(i)
```

```
    ENDDO
```

```
ENDDO
```

Notice: for $i=n$ we have:

DO $j=n+1,n$ which will not enter the loop

Assumed-size arrays

- The size of arrays passed to subprograms can be passed along with the array (FORTRAN 77 style):

Example:

```
SUBROUTINE SUB(A,lda,ni,nj)
```

```
INTEGER, INTENT(IN) :: lda ! Leading dimension(s)
```

```
INTEGER, INTENT(IN) :: ni,nj
```

```
REAL, DIMENSION(lda,*) :: A ! Last dimension can be left as: *
```

A = 0 ! Is illegal – since A is not assumed-shape (:,:) style

```
DO j=1,nj
```

```
    DO i=1,ni
```

```
        A(i,j) = ... ! This is legal
```

Assumed-size arrays

```
PROGRAM main
```

```
REAL, DIMENSION(7,5) :: field
```

```
CALL sub(field(3,4),7,3,2)
```

... and in the subroutine:

```
SUBROUTINE sub(A,lda,ni,nj)
```

```
INTEGER :: lda,ni,nj
```

```
REAL, DIMENSION(lda,*) :: A
```

```
DO j=1,nj
```

```
  DO i=1,ni
```

```
    A(i,j) = ...
```

! will access the elements: field(3:5,4:5)

Assumed-shape arrays

- For assumed-shape arrays (:-notation) the size of the arrays passed to subprogram can be inquired by the subprogram (but requires an INTERFACE).

- Example:

```
program main
```

```
real, dimension(300,300) :: matrix
```

```
interface
```

```
    subroutine sub(array)
```

```
        real, dimension(:, :) :: array
```

```
    end subroutine sub
```

```
end interface
```

```
call sub(matrix)
```

```
...
```

- The interface only passes the SHAPE – not the lower and upper bounds. Thus if the *actual* argument has the DIMENSION(-1,10) the *dummy* argument will be DIMENSION(1:12).

Assumed-shape arrays

Example:

```
SUBROUTINE SUB(field)
```

```
REAL, DIMENSION(:,:) [,POINTER] :: field
```

```
PRINT*, 'LBOUND(field) = ', LBOUND(field)
```

```
PRINT*, 'UBOUND(field) = ', UBOUND(field)
```

- A POINTER array is NOT considered assumed-shape, and if the actual and dummy arguments are given the POINTER or ALLOCATABLE attribute the print will return the bounds of the actual argument.
- If the dummy argument is an assumed-shape (non-POINTER/ALLOCATABLE) the LBOUND will return 1,1 and the UBOUND the total extend of the array.

Assumed-shape arrays

```
program main
  implicit none
  real, dimension(-1:10) :: data
  integer :: i
  interface
    subroutine sub(vector)
      real, dimension(:) :: vector
    end subroutine sub
  end interface

  call sub(data)
  do i=-1,4
    print*, 'data(i) = ', i, data(i)
  enddo
end program main

subroutine sub(vector)
  implicit none
  real, dimension(:) :: vector
  integer :: i

  do i=lbound(vector,1), ubound(vector,1)
    vector(i) = i
  enddo
  return
end subroutine sub
```

Will print:

```
-1 1.00000
0 2.00000
1 3.00000
2 4.00000
3 5.00000
4 6.00000
```

Automatic arrays

- Arrays can be declared in subprograms based on dummy arguments (so-called *automatic* arrays):

Example:

```
SUBROUTINE swap(a,b)
```

```
REAL, DIMENSION(:), INTENT(INOUT) :: a,b
```

```
REAL, DIMENSION(SIZE(a)) :: work ! Will be allocated on the stack !
```

```
REAL, DIMENSION(100000) :: aux ! Will be allocated on the stack !
```

```
work = a; a = b; b = work
```

```
END SUBROUTINE swap
```

DO NOT USE AUTOMATIC ARRAYS!

It may/will crash during:

- 1) The declaration of the work array**
- 2) or/and during the assignment**

Stack

- Automatic arrays are allocated on the *stack* – **without** a return / error status.
- Local arrays are (normally) allocated on the stack:
 - Local static arrays has no return status.
 - ALLOCATE of local POINTERS/ALLOCABLEs returns a status.
- The stack is usually of limited size (try the UNIX command: `ulimit -a`; gbar: ca. 10MB)
- Advice: place the subprogram in a module and place the local arrays in the module holding the subprogram – then stored on the heap.

Local arrays

Example: local arrays allocated on the stack:

```
SUBROUTINE swap(a,b)
```

```
REAL, DIMENSION(:), POINTER :: work ! Local array
```

```
REAL, DIMENSION(:), INTENT(INOUT) :: a,b
```

```
! PRO: ALLOCATE will at give return status;
```

```
! CON: the work array will/may be allocated on the (small)
```

```
! stack and easily fail !
```

```
ALLOCATE(work(SIZE(a)),STAT=info)
```

```
work = a; a = b; b = work
```

```
DEALLOCATE(work,STAT=info)
```

```
END SUBROUTINE swap
```

Local arrays

Example: always place local arrays into the module!

```
MODULE m_swap
```

```
REAL, DIMENSION(:), POINTER :: work
```

```
CONTAINS
```

```
  SUBROUTINE swap(a,b)
```

```
    REAL, DIMENSION(:), INTENT(INOUT) :: a,b
```

```
    ALLOCATE(work(SIZE(a)),STAT=info)
```

```
    work = a; a = b; b = work
```

```
    DEALLOCATE(work,STAT=info)
```

```
  END SUBROUTINE swap
```

```
END MODULE m_swap
```

Will crash if `size(a).NE.size(b)`
So: check the size of a and b
and decide what to do !
then use DO loops for the copy !
You can also keep the work array
allocated to save time if you call
the swap routine repeatedly.

Allocatable arrays

- The *rank* of an allocatable array is specified when it is defined.
- The *bounds* of the array are undefined until the ALLOCATE statement:

Example:

```
MODULE m_swap
```

```
    REAL, DIMENSION(:), ALLOCATABLE :: work
```

```
CONTAINS
```

```
    SUBROUTINE swap(a,b)
```

```
        REAL, DIMENSION(:), INTENT(INOUT) :: a,b
```

```
        ALLOCATE(work(SIZE(a)),STAT=info)
```

Re-allocatable arrays

- ALLOCATABLE (or POINTER) arrays cannot be *reallocated* (grow/shrink while preserving the content of the array) using a single call.
- To reallocate an array (A) do:
 - create a work array (W) of the size of A.
 - copy the data from A to W.
 - deallocate A.
 - allocate A with the new required size.
 - copy the data from the W to A.
 - deallocate W.

Re-allocatable arrays

- OR: to reallocate an (POINTER) array (A) do:
 - create work array (W) of the new desired size of A.
 - copy the data from A to W.
 - deallocate A.
 - let A point to W (saving the last copy back !)
- OR: use MOVE_ALLOC() for allocatable arrays:
 - create work array (W) of new desired size of A.
 - copy the data from A to W.
 - call move_alloc(W,A)

The ALLOCATE statement

- ALLOCATE(allocate-list,[STAT=stat])
 - allocate-list: allocate-object[(array-bounds-list)]
 - array-bounds-list: [lower-bounds:]upper-bounds

Example:

```
REAL, DIMENSION(:,:), ALLOCATABLE :: array1
```

```
REAL, DIMENSION(:,:), POINTER :: array2
```

```
INTEGER :: istat
```

```
ALLOCATE(array1(-10:5,5),STAT=istat)
```

```
ALLOCATE(array2(100,100),STAT=istat)
```

The ALLOCATE statement

- The allocate-object can be an ALLOCATABLE or POINTER array.
- Each lower/upper-bound is a scalar integer expression.
- The default lower-bound is 1.
- The number of array bounds in a list must match the rank of allocate-object.
- The STAT returns zero on success.
- If the STAT is absent and the allocation is unsuccessful, the program execution will stop.
- **ADVICE: ALWAYS use STAT and ALWAYS check the value of STAT !**

The ALLOCATE statement

- During the call to ALLOCATE the array boundary is undefined, thus:
 - `ALLOCATE(a(SIZE(b)),b(size(a)))` ! Is illegal
 - `ALLOCATE(a(n),b(SIZE(a)))` ! Is illegal
 - `ALLOCATE(a(n)); ALLOCATE(b(SIZE(a)))` ! Is legal
- To allocate an already allocated
 - ALLOCATABLE array is illegal.
 - POINTER array is legal (but leave the previous target inaccessible!).

The DEALLOCATE statement

- DEALLOCATE(allocate-object-list[,STAT=stat])

Example:

```
REAL, DIMENSION(:,:), ALLOCATABLE :: array1
```

```
REAL, DIMENSION(:,:), POINTER :: array2
```

```
INTEGER :: istat
```

```
DEALLOCATE(array1,STAT=istat)
```

```
DEALLOCATE(array2,STAT=istat)
```

```
! Or DEALLOCATE(array1,array2,STAT=istat)
```

The DEALLOCATE statement

- The STAT returns zero on success.
- If the STAT is absent and the deallocation is unsuccessful, the program execution will stop.
- **ADVICE: ALWAYS use STAT and ALWAYS check the value of STAT !**
- There must be no dependencies in the list of objects to allow a one-by-one deallocation.
- If an array is deallocated **other** pointers to the array are left *dangling* !

The NULLIFY statement

- NULLIFY(pointer-object-list)
 - Disassociates a pointer from its target.
 - Null pointer is different than undefined (dangling) and nullified pointers can be tested by the intrinsic routine: ASSOCIATED().
 - Nullify does not deallocate the target (if deallocation is needed DEALLOCATE should be used).

The WHERE statement

- The WHERE statement performs array operations for certain elements only:

[name:] WHERE (*logical-array-expr*)

array-variable=expr

[ELSE WHERE]

array-variable=expr

END WHERE [name]

Example:

WHERE ($a > 0.0$) $a = 1.0/a$! a is a real matrix

DO NOT USE IT!

- 1) It is NOT more efficient than the DO.
- 2) It required the SHAPE to be identical for all arrays.
- 3) It may allocate (hidden) extra memory! (and crash in the process).

The FORALL statement

- FORALL is like DO but does not require order:

```
[name:] FORALL (index=lower:upper[:stride] &  
    [index=lower:upper[:stride] ...  
    [scalar-local-expr])  
    [body]
```

```
END FORALL [name]
```

Example:

```
FORALL (i=1:N)
```

```
    A(i) = 1.0
```

```
END FORALL
```

Example:

```
FORALL (i=1:n)
```

```
    WHERE (a(i,:) == 0) a(i,:)=i
```

```
    b(i,:)=i/a(i,:)
```

```
END FORALL
```


PURE procedures

- A procedure (subroutine/function) is PURE if:
 - A *function* does not alter any dummy arguments.
 - It does not alter any part of a variable accessed by host or use association (through module data).
 - It contains no local variables with the SAVE attribute.
 - Performs no operation on an external file.
 - Contains no stop statement.
- The compiler will produce an error if these rules are violated.

PURE procedures

- To assert that a procedure has no side-effects add the PURE keyword to the subroutine or function statement.

Example:

```
PURE FUNCTION distance(p,q)
```

- An interface block is required if a pure procedure is to be used as pure (i.e. pure procedures can also be used without an interface).

ELEMENTAL procedures

- Intrinsic functions are *elemental* – i.e. the result is given the shape of the input.

Example:

REAL :: scalar

REAL, DIMENSION(100) :: vector

scalar = 1.0; scalar = SQRT(scalar)

vector(1:100) = 3.0; vector = SQRT(vector)

ELEMENTAL procedures

- Non-intrinsic procedures can be *elemental* if:
 1. The declaration contains the prefix ELEMENTAL.
 2. All the arguments are conformable (same shape).
 3. The procedure is PURE.
 4. All dummy arguments and function results must be **scalar** variables without the pointer attribute.

Example:

```
ELEMENTAL SUBROUTINE swap(a,b)
```

```
REAL, INTENT(INOUT) :: a,b
```

```
REAL :: work
```

```
work = a; a = b; b = work
```

```
END SUBROUTINE swap
```

This swap routine can now be called with any REAL scalar or array data !

ELEMENTAL procedures

```
PROGRAM main
INTERFACE swap
  ELEMENTAL SUBROUTINE swap(a,b)
    REAL, INTENT(INOUT) :: a,b
  END SUBROUTINE swap
END INTERFACE swap
REAL                :: sx,sy
REAL, DIMENSION(2) :: vx,vy
sx = 1.0; sy = 3.0
vx(1) = 1.0; vx(2) = 2.0
vy(1) = -23.; vy(2) = -345
CALL swap(sx,sy) ! We can call swap() with any rank !
CALL swap(vx,vy) ! But we need overloading for different KIND !
END PROGRAM main
```

ELEMENTAL procedures are
EXTREMELY powerful !!
AGAIN: be careful: lots of things
are going on BEHIND the SCENE
and ELEMENTAL functions will be
SLOWER than ordinary overloaded
functions

Explicit overloading

```
MODULE oswap
  REAL :: swork ! Work array for scalar swap
  REAL, DIMENSION(:), ALLOCATABLE :: vwork ! Work array for vector swap
  INTERFACE swap
    MODULE PROCEDURE sswap, vswap
  END INTERFACE swap
  CONTAINS
    SUBROUTINE sswap(a,b) ! Scalar version
      REAL :: a,b
      swork = a; a = b; b = swork
    END SUBROUTINE sswap
    SUBROUTINE vswap(a,b) ! Vector version
      REAL, DIMENSION(:) :: a,b
      ALLOCATE(vwork(SIZE(a)))
      vwork = a; a = b; b = vwork
      DEALLOCATE(vwork)
    END SUBROUTINE vswap
END MODULE oswap
```

A safe alternative to ELEMENTAL
procedures:
use explicit overloading !

ELEMENTAL procedures

- If a generic procedure reference is consistent with BOTH an elemental and a non-elemental procedure, the non-elemental procedure is invoked.
- Thus we can write efficient specific versions and leave the ELEMENTAL to the general case!

Array sub-objects

- Subsections of an array may be extracted by:
 $a(i,1:n)$! Elements 1 to n of row i
 $a(1:m,j)$! Elements 1 to m of column j
 $a(i,:)$! The whole row i
 $a(i,1:n:3)$! Elements 1,4,... of row i
 $v([1,7,3,2])$! Is a vector of length 4 with the elements: $v(1)$, $v(7)$,
 $v(3)$, $v(2)$

Array sub-objects

equivalently:

$v(\text{list})$, where

$\text{list}(1) = 1; \text{list}(2) = 7$

$\text{list}(3) = 3; \text{list}(4) = 2$

- Vector subscript must be unique when appearing on the LHS:

$v(/1,7,3,7/) = (/1,2,3,4/) !$ Is illegal

since 2 and 4 cannot both be stored in 7

Array sub-objects

- Subscript: lower:upper:stride
 - Lower and upper may exceed the range of the array
 - Stride has to be non-zero.

Example:

$A(i, 1:11:3)$! Address 1,4,7,10 of column i

$A(i, 10:1:-3)$! Address 10,7,4,1 of column i

Arrays of pointers

- Array of pointers does not exist in FORTRAN.
- The same effect can be achieved by an array of user defined type:

Example: lower-triangular matrix

```
TYPE row
```

```
    REAL, DIMENSION(:), POINTER :: r
```

```
END TYPE row
```

```
TYPE (row), DIMENSION(n):: s,t
```

```
DO i=1,n
```

```
    ALLOCATE(t(i)%r(i),STAT=istat)
```

```
ENDDO
```

Pointers as aliases

- Example:

```
REAL, DIMENSION(:,:), POINTER :: window
```

```
window => table(m:n,p:q)
```

```
DO j=1,q-p+1
```

```
    DO i=1,n-m+1
```

```
        window(i,j) = ...
```

```
    ENDDO
```

```
ENDDO
```